

### CONSTRAINT SATISFACTION

A constraint satisfaction problem (CSP) is defined by:

$$\begin{aligned} \text{variables } X &= \{X_1, X_2, \dots\} \\ \text{constraints } C &= \{C_1, C_2, \dots\} \end{aligned}$$

- Each *variable*  $X_i$  has a non-empty domain  $D_i$  of possible values.
- Each *constraint*  $C_i$  involves some subset of the variables—the *scope* of the constraint—and specifies the allowable combinations of values for that subset.
- An *assignment* that does not violate any constraints is called *consistent* (or solution).

### RATIONAL METAREASONING

- A problem-solving agent can perform *base-level* actions from a known set  $\{A_i\}$ .
- Before committing to an action, the agent may perform a sequence of *meta-level* deliberation actions from a set  $\{S_j\}$ .
- At any given time there is a base-level action  $A_\alpha$  that maximizes the agent's *expected utility*.

The **net VOI**  $V(S_j)$  of action  $S_j$  is the intrinsic VOI  $\Lambda_j$  less the cost of  $S_j$ :

$$V(S_j) = \Lambda(S_j) - C(S_j)$$

The **intrinsic VOI**  $\Lambda(S_j)$  is the expected difference between the intrinsic expected utilities of the new and the old selected base-level action, computed after the meta-level action is taken:

$$\Lambda(S_j) = E(E(U(A_\alpha^j)) - E(U(A_\alpha)))$$

- $S_{j_{\max}}$  that maximizes the net VOI is performed:  
 $j_{\max} = \arg \max_j V(S_j)$   
if  $V(S_{j_{\max}}) > 0$
- Otherwise,  $A_\alpha$  is performed.

### ACKNOWLEDGMENTS

- IMG4 Consortium under the MAGNET program of the Israeli Ministry of Trade and Industry
- Israel Science Foundation grant 305/09
- Lynne and William Frankel Center for Computer Sciences
- Paul Ivanier Center for Robotics Research and Production Management

### CONCLUSIONS

- This work suggests a model for adaptive deployment of value ordering heuristics in algorithms for constraint satisfaction problems.
- As a case study, the model was applied to a value-ordering heuristic based on solution count estimates, and a steady improvement was achieved compared to always computing the estimates.
- For many problem instances the optimum performance is achieved when solution counts are estimated only in a small number of search states.

### OVERVIEW

Heuristics are crucial tools in decreasing search effort in various areas of AI. A heuristic must provide useful information to the search algorithm, and be efficient to compute.

Overhead of some well-known heuristics may outweigh the gain.

Such heuristics should be deployed selectively, based on principles of rational metareasoning.

### Case Study

- CSP backtracking search algorithms typically employ variable-ordering and value-ordering heuristics.
- Many value ordering heuristics are computationally heavy, e.g. heuristics based on solution count estimates.
- Principles of rational metareasoning can be applied to decide when to deploy the heuristics.

### VALUE ORDERING

Value ordering heuristics provide information about:

- $T_i$ —the expected time to find a solution containing an assignment  $X_k = y_{ki}$ ;
- $p_i$ —the probability that there is no solution consistent with  $X_k = y_{ki}$ .

The expected remaining search time in the subtree under  $X_k$  for ordering  $\omega$  is given by:

$$T^{\omega} = T_{\omega(1)} + \sum_{i=2}^{|D_k|} T_{\omega(i)} \prod_{j=1}^{i-1} p_{\omega(j)}$$

- The current optimal base-level action is picking the  $\omega$  which optimizes  $T^{\omega}$ .  $T^{\omega}$  is minimal if the values are sorted by increasing order of  $\frac{T_i}{1-p_i}$ .
- The intrinsic VOI  $\Lambda_i$  of estimating  $T_i$ ,  $p_i$  for the  $i$ th assignment is the expected decrease in the expected search time:  $\Lambda_i = E[T^{\omega_{-}} - T^{\omega_{+}i}]$ .
- Computing new estimates (with overhead  $T^c$ ) for values  $T_i$ ,  $p_i$  is beneficial just when the net VOI is positive:  $V_i = \Lambda_i - T^c$ .

### MAIN RESULTS

#### Rational Value Ordering

The intrinsic VOI  $\Lambda_i$  of invoking the heuristic can be approximated as:

$$\Lambda_i \approx E[(T_1 - T_i)|D_k| \mid T_i < T_1]$$

#### VOI of Solution Count Estimates

The net VOI  $V$  of estimating a solution count can be approximated as:

$$V \propto |D_k| e^{-\nu} \sum_{n=n_{\max}}^{\infty} \left( \frac{1}{n_{\max}} - \frac{1}{n} \right) \frac{\nu^n}{n!} - \gamma$$

where the constant  $\gamma$  depends on the search algorithm and the heuristic, rather than on the CSP instance, and can be learned offline. The infinite sum  $\sum_{n=n_{\max}}^{\infty} \left( \frac{1}{n_{\max}} - \frac{1}{n} \right) \frac{\nu^n}{n!}$  is rapidly converging and can be computed efficiently.

### ALGORITHM

#### SC-based Rational Value Ordering

```

procedure VALUEORDERING-SC( $csp, X_k, N$ )
   $D \leftarrow D_k$ ,  $n_{\max} \leftarrow \frac{N}{|D|}$ 
  for all  $i$  in  $1..|D|$  do  $n_i \leftarrow n_{\max}$ 
  while  $V(n_{\max}) > 0$ 
    choose  $y_{ki} \in D$  arbitrarily
     $D \leftarrow D \setminus \{y_{ki}\}$ 
     $csp' \leftarrow csp$  with  $D_k = \{y_{ki}\}$ 
     $n_i \leftarrow \text{ESTIMATESOLUTIONCOUNT}(csp')$ 
    if  $n_i > n_{\max}$  then  $n_{\max} \leftarrow n_i$ 
  end while
   $D_{ord} \leftarrow \text{sort } D_k$ 
    by non-increasing  $n_i$ 
  return  $D_{ord}$ 
end procedure

```

### EXPERIMENTS

#### Benchmarks

CSP benchmarks from CSP Solver Competition 2005 were used. 14 benchmarks were solved for  $\gamma = 0$  and the exponential range  $\gamma \in \{10^{-7}, 10^{-6}, \dots, 1\}$ , as well as with the minimum-conflicts heuristic and the pAC heuristic.

The maximum improvement is achieved when the solution count is estimated only in a small fraction of occasions selected using rational metareasoning.

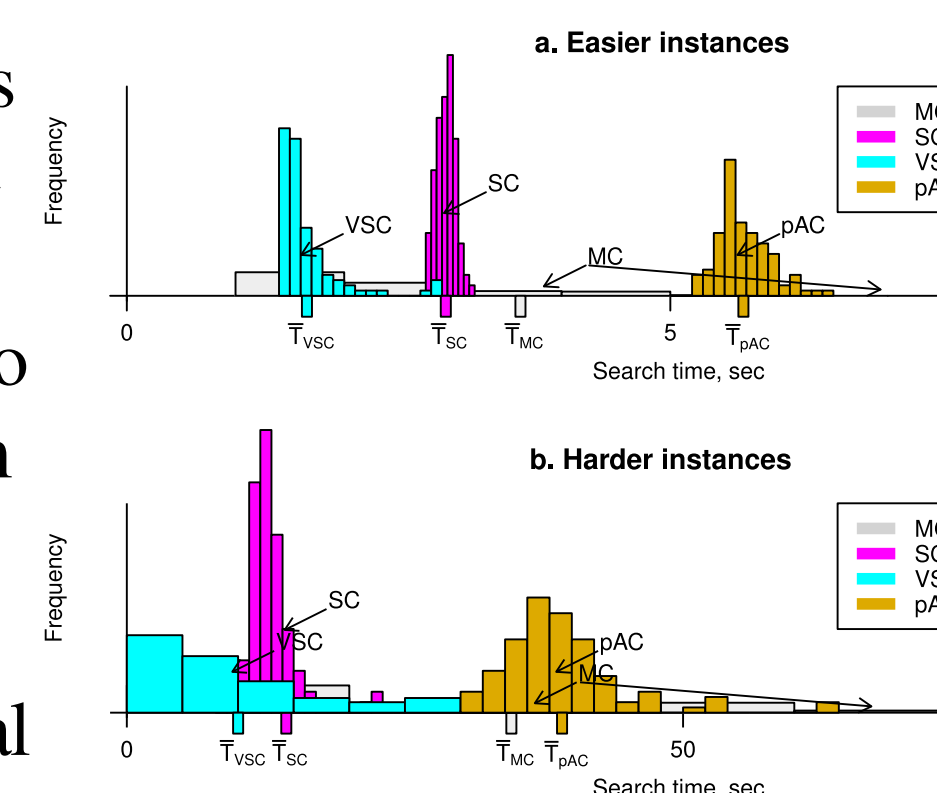
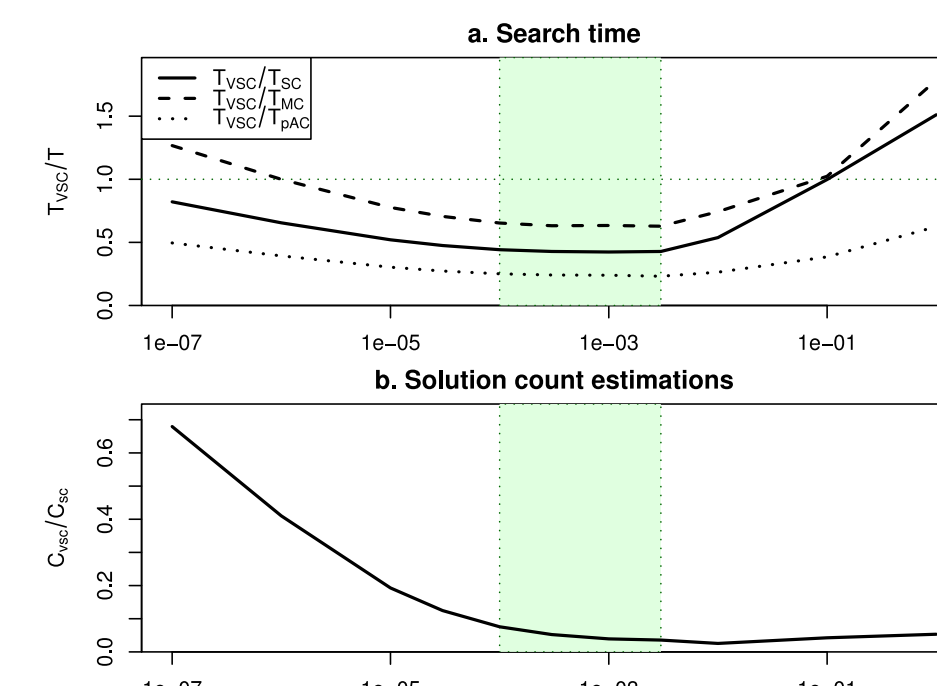
#### Random instances

Based on the results on benchmarks, we chose  $\gamma = 10^{-3}$ , and applied it to two sets of 100 problem instances. Exhaustive deployment, rational deployment, the minimum conflicts heuristic, and the pAC heuristic were compared.

The value of  $\gamma$  chosen based on a small set of hard instances gave good results on a set of instances with different parameters and of varying hardness.

#### Generalized Sudoku

- Real-world problem instances often have much more structure than random instances generated according to Model RB.
- We repeated the experiments on randomly generated Generalized Sudoku instances—a highly structured domain.
- Relative performance on Generalized Sudoku was similar to Model RB.



### FUTURE WORK

- Generalization of the VOI to deploy different types of heuristics for CSP.
- Explicit evaluation of the quality of the distribution model, coupled with a better candidate model of the distribution.
- Application to search in other domains, especially to heuristics for planning; in particular, examining whether the meta-reasoning scheme can improve reasoning over deployment based solely on learning.

